

Esper as DAQ expert system backbone - evaluation

DAQ Topical

Tomasz Bawej, supervisor: Hannes Sakulin

20.02.2014

1 Introduction to Complex Event Processing with Esper

- Complex Event Processing
- Overview of Esper

2 Esper Event Processing Language by example

- Streams and windows
- Contexts
- Patterns
- Match recognize: detect state changes of DAQ

3 Expert system prototype

- Prototype overview
- Consistency with DaqDoctor
- Performance

4 Evaluation of Esper

5 Discussion

Goals

Derive useful information from streams of data and respond to it as quickly as possible.

Goals

Derive useful information from streams of data and respond to it as quickly as possible.

Means

- ▶ Event-pattern detection
- ▶ Event abstraction, filtering, aggregation, transformation and correlation
- ▶ Modeling event hierarchies and abstracting event-driven processes

What is Esper?

- ▶ Event series analysis and event correlation engine.
- ▶ Lightweight kernel written in Java (plus a version for .NET).
- ▶ Available both as open source and an Enterprise Edition.

What is Esper?

- ▶ Event series analysis and event correlation engine.
- ▶ Lightweight kernel written in Java (plus a version for .NET).
- ▶ Available both as open source and an Enterprise Edition.

Notable customers:

- ▶ PayPal
- ▶ Huawei
- ▶ Oracle

- ▶ *Events* pass through the Esper engine in *streams*.

Event definitions

Statement registration must be preceded by defining all event types the statement refers to.

The event can be represented by an instance of:

- ▶ `java.util.Map`
- ▶ `java.lang.Object[]`
- ▶ POJO
- ▶ XML document

- ▶ *Events* pass through the Esper engine in *streams*.
- ▶ *Windows* utilise *views* to retain events.

Event definitions

Statement registration must be preceded by defining all event types the statement refers to.

The event can be represented by an instance of:

- ▶ `java.util.Map`
- ▶ `java.lang.Object[]`
- ▶ POJO
- ▶ XML document

- ▶ *Events* pass through the Esper engine in *streams*.
- ▶ *Windows* utilise *views* to retain events.
- ▶ *Views* can be of four types:

Event definitions

Statement registration must be preceded by defining all event types the statement refers to.

The event can be represented by an instance of:

- ▶ `java.util.Map`
- ▶ `java.lang.Object[]`
- ▶ POJO
- ▶ XML document

- ▶ *Events* pass through the Esper engine in *streams*.
- ▶ *Windows* utilise *views* to retain events.
- ▶ *Views* can be of four types:
 1. *data win*: views - *n* (seconds of) events, expiry expression, batch, accumulating...
 2. *std*: views - unique, grouped sub-views, size, first, last
 3. *stat*: views - statistics for specified properties
 4. *ext*: views - sorting

Event definitions

Statement registration must be preceded by defining all event types the statement refers to.

The event can be represented by an instance of:

- ▶ `java.util.Map`
- ▶ `java.lang.Object[]`
- ▶ POJO
- ▶ XML document

- ▶ *Events* pass through the Esper engine in *streams*.
- ▶ *Windows* utilise *views* to retain events.
- ▶ *Views* can be of four types:
 1. *data win*: views - *n* (seconds of) events, expiry expression, batch, accumulating...
 2. *std*: views - unique, grouped sub-views, size, first, last
 3. *stat*: views - statistics for specified properties
 4. *ext*: views - sorting
- ▶ *Statements* are instructions for the engine, executed whenever applicable.
 - ▶ Statements can refer to events entering the window, leaving it or both. (*istream*)

Event definitions

Statement registration must be preceded by defining all event types the statement refers to.

The event can be represented by an instance of:

- ▶ `java.util.Map`
- ▶ `java.lang.Object[]`
- ▶ POJO
- ▶ XML document

- ▶ *Events* pass through the Esper engine in *streams*.
- ▶ *Windows* utilise *views* to retain events.
- ▶ *Views* can be of four types:
 1. *data win*: views - *n* (seconds of) events, expiry expression, batch, accumulating...
 2. *std*: views - unique, grouped sub-views, size, first, last
 3. *stat*: views - statistics for specified properties
 4. *ext*: views - sorting
- ▶ *Statements* are instructions for the engine, executed whenever applicable.
 - ▶ Statements can refer to events entering the window, leaving it or both. (*irstream*)
- ▶ *Listeners* (*subscribers*) can be registered with *statements* to receive their updates.

Event definitions

Statement registration must be preceded by defining all event types the statement refers to.

The event can be represented by an instance of:

- ▶ `java.util.Map`
- ▶ `java.lang.Object[]`
- ▶ POJO
- ▶ XML document

Statements can be constructed using an SQL-like Event Processing Language (EPL).

SQL features present in EPL

- ▶ tables, rows and columns translate to windows, events and properties
- ▶ clauses (e.g. CREATE, SELECT, INSERT, DELETE, UPDATE, MERGE, FROM, WHERE, GROUP BY, HAVING and ORDER BY)
- ▶ aggregation functions (e.g. `sum()`, `avg()`, `count()`, `max()`)
- ▶ triggers resembled by event driven queries (ON...)

Statements can be constructed using an SQL-like Event Processing Language (EPL).

SQL features present in EPL

- ▶ tables, rows and columns translate to windows, events and properties
- ▶ clauses (e.g. CREATE, SELECT, INSERT, DELETE, UPDATE, MERGE, FROM, WHERE, GROUP BY, HAVING and ORDER BY)
- ▶ aggregation functions (e.g. sum(), avg(), count(), max())
- ▶ triggers resembled by event driven queries (ON...)

Simple SQL-like esper query.

```
select systime, srctime, fromState, toState from  
    DaqStateChangeStream;
```

Streams and windows: prompt on average stream A rate $> 500\text{Hz}$

```
create objectarray schema StreamRates as (name String, rate double,  
      srctime long);
```

Streams and windows: prompt on average stream A rate $> 500\text{Hz}$

Stream definition, populated by summing data over Storage Managers.

```
create objectarray schema StreamRates as (name String, rate double,  
      srctime long);
```


Streams and windows: prompt on average stream A rate $> 500\text{Hz}$

Stream definition, populated by summing data over Storage Managers.

```
create objectarray schema StreamRates as (name String, rate double,  
      srctime long);
```

```
create window RatesWin.win:time(30 sec) as select * from StreamRates;
```

Streams and windows: prompt on average stream A rate $> 500\text{Hz}$

Stream definition, populated by summing data over Storage Managers.

```
create objectarray schema StreamRates as (name String, rate double,  
      srctime long);
```

Window retaining events for 30 seconds. Copies structure from RatesStream.

```
create window RatesWin.win:time(30 sec) as select * from StreamRates;
```

Streams and windows: prompt on average stream A rate $> 500\text{Hz}$

Stream definition, populated by summing data over Storage Managers.

```
create objectarray schema StreamRates as (name String, rate double,  
      srctime long);
```

Window retaining events for 30 seconds. Copies structure from RatesStream.

```
create window RatesWin.win:time(30 sec) as select * from StreamRates;
```

```
insert into RatesWin select * from StreamRates(name="A");
```

Streams and windows: prompt on average stream A rate $> 500\text{Hz}$

Stream definition, populated by summing data over Storage Managers.

```
create objectarray schema StreamRates as (name String, rate double,  
      srctime long);
```

Window retaining events for 30 seconds. Copies structure from RatesStream.

```
create window RatesWin.win:time(30 sec) as select * from StreamRates;
```

Forwards events from RatesStream into the window.

```
insert into RatesWin select * from StreamRates(name="A");
```

Streams and windows: prompt on average stream A rate > 500Hz

Stream definition, populated by summing data over Storage Managers.

```
create objectarray schema StreamRates as (name String, rate double,  
      srctime long);
```

Window retaining events for 30 seconds. Copies structure from RatesStream.

```
create window RatesWin.win:time(30 sec) as select * from StreamRates;
```

Forwards events from RatesStream into the window.

```
insert into RatesWin select * from StreamRates(name="A");
```

```
@Verbose(label="out", fields={"avgRate"}, extraNfo="A rate high")  
select avg(rate) as avgRate from RatesWin  
  having avg(rate)>500 output first every 10 seconds;
```

Streams and windows: prompt on average stream A rate > 500Hz

Stream definition, populated by summing data over Storage Managers.

```
create objectarray schema StreamRates as (name String, rate double,  
      srctime long);
```

Window retaining events for 30 seconds. Copies structure from RatesStream.

```
create window RatesWin.win:time(30 sec) as select * from StreamRates;
```

Forwards events from RatesStream into the window.

```
insert into RatesWin select * from StreamRates(name="A");
```

Java-style annotation.

```
@Verbose(label="out", fields={"avgRate"}, extraNfo="A rate high")  
select avg(rate) as avgRate from RatesWin  
  having avg(rate)>500 output first every 10 seconds;
```

Annotation - extra information about a statement

- ▶ Esper provides a set of built-in annotations.
- ▶ Custom annotations can also be provided and used.
 - ▶ @Verbose marks the statement as producing relevant output.
 - ▶ @Watched marks statements as producing continuous output.

From now on the @Verbose annotation will be omitted for the sake of clarity.

Streams and windows: prompt on average stream A rate > 500Hz

Stream definition, populated by summing data over Storage Managers.

```
create objectarray schema StreamRates as (name String, rate double,  
      srctime long);
```

Window retaining events for 30 seconds. Copies structure from RatesStream.

```
create window RatesWin.win:time(30 sec) as select * from StreamRates;
```

Forwards events from RatesStream into the window.

```
insert into RatesWin select * from StreamRates(name="A");
```

Java-style annotation.

```
@Verbose(label="out", fields={"avgRate"}, extraNfo="A rate high")  
select avg(rate) as avgRate from RatesWin  
  having avg(rate)>500 output first every 10 seconds;
```

Clause limiting the output rate.

Annotation - extra information about a statement

- ▶ Esper provides a set of built-in annotations.
- ▶ Custom annotations can also be provided and used.
 - ▶ @Verbose marks the statement as producing relevant output.
 - ▶ @Watched marks statements as producing continuous output.

From now on the @Verbose annotation will be omitted for the sake of clarity.

```
select average
  from StreamRates.win:time(30 seconds).stat:uni(rate)
 where average>500 output first every 10 seconds;
```


average is a field of the stat:uni (univariate) view.

```
select average
  from StreamRates.win:time(30 seconds).stat:uni(rate)
 where average>500 output first every 10 seconds;
```

Streams and windows: prompt on stream A rate $> 500\text{Hz}$

average is a field of the stat:uni (univariate) view.

Chained views: the second is applied to the results of applying the first.

```
select average
from StreamRates.win:time(30 seconds).stat:uni(rate)
where average>500 output first every 10 seconds;
```

Streams and windows: extend frlcontrollerLink with fedId

```
create objectarray schema FrlCtlLnk(fedSrcId int) copyfrom  
  frlcontrollerLink;
```

Streams and windows: extend frlcontrollerLink with fedId

Extending a copied schema definition.

```
create objectarray schema FrlCtlLnk(fedSrcId int) copyfrom  
  frlcontrollerLink;
```

Streams and windows: extend frlcontrollerLink with fedId

Extending a copied schema definition.

```
create objectarray schema FrlCtlLnk(fedSrcId int) copyfrom  
    frlcontrollerLink;
```

```
insert into FrlCtlLnk select *,  
    fedSrcId(x.context,x.slotNumber,x.linkNumber,CmsHw.FRL) as fedSrcId  
from frlcontrollerLink as x where  
fedSrcId(x.context,x.slotNumber,x.linkNumber,CmsHw.FRL)  
in (select fedSrcId from FedMask(slinkEnabled=true));
```

Streams and windows: extend frlcontrollerLink with fedId

Extending a copied schema definition.

```
create objectarray schema FrlCtlLnk(fedSrcId int) copyfrom  
  frlcontrollerLink;
```

Copying all the frlcontrollerLink fields.

```
insert into FrlCtlLnk select *,  
  fedSrcId(x.context,x.slotNumber,x.linkNumber,CmsHw.FRL) as fedSrcId  
from frlcontrollerLink as x where  
fedSrcId(x.context,x.slotNumber,x.linkNumber,CmsHw.FRL)  
in (select fedSrcId from FedMask(slinkEnabled=true));
```

Streams and windows: extend frlcontrollerLink with fedId

Extending a copied schema definition.

```
create objectarray schema FrlCtlLnk(fedSrcId int) copyfrom  
  frlcontrollerLink;
```

Inserting fedSrcId returned from a Java helper method.

Copying all the frlcontrollerLink fields.

```
insert into FrlCtlLnk select *,  
  fedSrcId(x.context,x.slotNumber,x.linkNumber,CmsHw.FRL) as fedSrcId  
from frlcontrollerLink as x where  
fedSrcId(x.context,x.slotNumber,x.linkNumber,CmsHw.FRL)  
in (select fedSrcId from FedMask(slinkEnabled=true));
```

Streams and windows: extend frlcontrollerLink with fedId

Extending a copied schema definition.

```
create objectarray schema FrlCtlLnk(fedSrcId int) copyfrom  
  frlcontrollerLink;
```

Inserting fedSrcId returned from a Java helper method.

Copying all the frlcontrollerLink fields.

```
insert into FrlCtlLnk select *,  
  fedSrcId(x.context,x.slotNumber,x.linkNumber,CmsHw.FRL) as fedSrcId  
from frlcontrollerLink as x where  
fedSrcId(x.context,x.slotNumber,x.linkNumber,CmsHw.FRL)  
in (select fedSrcId from FedMask(slinkEnabled=true));
```

Subquery against a named window.

Streams and windows: extend frlcontrollerLink with fedId

Extending a copied schema definition.

```
create objectarray schema FrlCtlLnk(fedSrcId int) copyfrom  
  frlcontrollerLink;
```

Inserting fedSrcId returned from a Java helper method.

Copying all the frlcontrollerLink fields.

```
insert into FrlCtlLnk select *,  
  fedSrcId(x.context,x.slotNumber,x.linkNumber,CmsHw.FRL) as fedSrcId  
from frlcontrollerLink as x where  
fedSrcId(x.context,x.slotNumber,x.linkNumber,CmsHw.FRL)  
in (select fedSrcId from FedMask(slinkEnabled=true));
```

Subquery against a named window.

An imported enum.

Contexts: computing FED backpressure

```
create objectarray schema FrlBackpressureStream as  
  (fedSrcId Integer, bpFraction double, timestamp long);
```

Contexts: computing FED backpressure

```
create objectarray schema FrlBackpressureStream as  
  (fedSrcId Integer, bpFraction double, timestamp long);
```

```
create context BpPerFedId  
  context RunOngoingNested start RunStart end RunStop,  
  context BpPerFedIdNested partition by fedSrcId from FrlCtlLnk;
```

Contexts: computing FED backpressure

```
create objectarray schema FrlBackpressureStream as  
  (fedSrcId Integer, bpFraction double, timestamp long);
```

A compound context - product of all sub-contexts.

```
create context BpPerFedId  
  context RunOngoingNested start RunStart end RunStop,  
  context BpPerFedIdNested partition by fedSrcId from FrlCtlLnk;
```

Contexts: computing FED backpressure

```
create objectarray schema FrlBackpressureStream as  
  (fedSrcId Integer, bpFraction double, timestamp long);
```

A compound context - product of all sub-contexts.

A non-overlapping sub-context.

```
create context BpPerFedId  
  context RunOngoingNested start RunStart end RunStop,  
  context BpPerFedIdNested partition by fedSrcId from FrlCtlLnk;
```

Contexts: computing FED backpressure

```
create objectarray schema FrlBackpressureStream as  
  (fedSrcId Integer, bpFraction double, timestamp long);
```

A compound context - product of all sub-contexts.

A non-overlapping sub-context.

```
create context BpPerFedId  
  context RunOngoingNested start RunStart end RunStop,  
  context BpPerFedIdNested partition by fedSrcId from FrlCtlLnk;
```

A sub-context partitioned by fedSrcId.

Contexts: computing FED backpressure

```
create objectarray schema FrlBackpressureStream as  
  (fedSrcId Integer, bpFraction double, timestamp long);
```

A compound context - product of all sub-contexts.

A non-overlapping sub-context.

```
create context BpPerFedId  
  context RunOngoingNested start RunStart end RunStop,  
  context BpPerFedIdNested partition by fedSrcId from FrlCtlLnk;
```

A sub-context partitioned by fedSrcId.

```
context BpPerFedId  
  insert into  
    FrlBackpressureStream  
  select  
    fedSrcId,  
    timestamp.getTime() as timestamp,  
    (fifoAlmostFullCnt - prior(1, fifoAlmostFullCnt))  
      / ((clockCount - prior(1, clockCount)) as bpFraction  
  from  
    FrlCtlLnk  
  where  
    clockCount > prior(1, clockCount);
```

Contexts: computing FED backpressure

```
create objectarray schema FrlBackpressureStream as  
  (fedSrcId Integer, bpFraction double, timestamp long);
```

A compound context - product of all sub-contexts.

A non-overlapping sub-context.

```
create context BpPerFedId  
  context RunOngoingNested start RunStart end RunStop,  
  context BpPerFedIdNested partition by fedSrcId from FrlCtlLnk;
```

A sub-context partitioned by fedSrcId.

```
context BpPerFedId  
  insert into  
    FrlBackpressureStream  
  select  
    fedSrcId,  
    timestamp.getTime() as timestamp,  
    (fifoAlmostFullCnt - prior(1, fifoAlmostFullCnt))  
      / ((clockCount - prior(1, clockCount)) as bpFraction  
  from  
    FrlCtlLnk  
  where  
    clockCount > prior(1, clockCount);
```

Compute backpressure value.

Contexts: computing FED backpressure

```
create objectarray schema FrlBackpressureStream as  
  (fedSrcId Integer, bpFraction double, timestamp long);
```

A compound context - product of all sub-contexts.

A non-overlapping sub-context.

```
create context BpPerFedId  
  context RunOngoingNested start RunStart end RunStop,  
  context BpPerFedIdNested partition by fedSrcId from FrlCtlLnk;
```

A sub-context partitioned by fedSrcId.

```
context BpPerFedId  
  insert into  
    FrlBackpressureStream  
  select  
    fedSrcId,  
    timestamp.getTime() as timestamp,  
    (fifoAlmostFullCnt - prior(1, fifoAlmostFullCnt))  
      / ((clockCount - prior(1, clockCount)) as bpFraction  
  from  
    FrlCtlLnk  
  where  
    clockCount > prior(1, clockCount);
```

Compute backpressure value.

prior single row function,
returns i-th previous event in order of arrival

Patterns: notify if jobcontrol flashlist is not updated within a minute

```
create variable int timeout = 50;
```

Patterns: notify if jobcontrol flashlist is not updated within a minute

Variable holding the timeout value minus update interval.

```
create variable int timeout = 50;
```

Patterns: notify if jobcontrol flashlist is not updated within a minute

Variable holding the timeout value minus update interval.

```
create variable int timeout = 50;
```

```
create context jobctlHostOutdated
context ByHost partition by hostname from jobcontrol,
context Outdated start pattern[a=jobcontrol->
    (timer:interval(timeout) and not jobcontrol(hostname=a.hostname))]
end pattern [jobcontrol(hostname=a.hostname)];
```

Patterns: notify if jobcontrol flashlist is not updated within a minute

Variable holding the timeout value minus update interval.

```
create variable int timeout = 50;
```

Creates a separate context partition per each hostname.

```
create context jobctlHostOutdated  
context ByHost partition by hostname from jobcontrol,  
context Outdated start pattern[a=jobcontrol->  
    (timer:interval(timeout) and not jobcontrol(hostname=a.hostname))]  
end pattern [jobcontrol(hostname=a.hostname)];
```

Patterns: notify if jobcontrol flashlist is not updated within a minute

Variable holding the timeout value minus update interval.

```
create variable int timeout = 50;
```

Creates a separate context partition per each hostname.

```
create context jobctlHostOutdated  
context ByHost partition by hostname from jobcontrol,  
context Outdated start pattern[a=jobcontrol->  
    (timer:interval(timeout) and not jobcontrol(hostname=a.hostname))]  
end pattern [jobcontrol(hostname=a.hostname)];
```

Starts when the pattern detects
an absence of flashlist.

Patterns: notify if jobcontrol flashlist is not updated within a minute

Variable holding the timeout value minus update interval.

```
create variable int timeout = 50;
```

Creates a separate context partition per each hostname.

```
create context jobctlHostOutdated  
context ByHost partition by hostname from jobcontrol,  
context Outdated start pattern[a=jobcontrol->  
    (timer:interval(timeout) and not jobcontrol(hostname=a.hostname))]  
end pattern [jobcontrol(hostname=a.hostname)];
```

Ends when the flashlist
with the same hostname arrives.

Starts when the pattern detects
an absence of flashlist.

Patterns: notify if jobcontrol flashlist is not updated within a minute

Variable holding the timeout value minus update interval.

```
create variable int timeout = 50;
```

Creates a separate context partition per each hostname.

```
create context jobctlHostOutdated
context ByHost partition by hostname from jobcontrol,
context Outdated start pattern[a=jobcontrol->
(timer:interval(timeout) and not jobcontrol(hostname=a.hostname))]
end pattern [jobcontrol(hostname=a.hostname)];
```

Ends when the flashlist
with the same hostname arrives.

Starts when the pattern detects
an absence of flashlist.

Accessing context property.

```
context jobctlHostOutdated
select context.ByHost.key1 as hostname,
date(current_timestamp()) as systime,
(current_timestamp()-context.Outdated.a.fetchstamp)/1000.0
as secondsSinceUpdate
from pattern [every timer:interval(10 seconds)];
```


Patterns: notify if jobcontrol flashlist is not updated within a minute

Variable holding the timeout value minus update interval.

```
create variable int timeout = 50;
```

Creates a separate context partition per each hostname.

```
create context jobctlHostOutdated
context ByHost partition by hostname from jobcontrol,
context Outdated start pattern[a=jobcontrol->
(timer:interval(timeout) and not jobcontrol(hostname=a.hostname))]
end pattern [jobcontrol(hostname=a.hostname)];
```

Ends when the flashlist
with the same hostname arrives.

Starts when the pattern detects
an absence of flashlist.

Accessing context property.

```
context jobctlHostOutdated
select context.ByHost.key1 as hostname,
date(current_timestamp()) as systime,
(current_timestamp()-context.Outdated.a.fetchstamp)/1000.0
as secondsSinceUpdate
from pattern [every timer:interval(10 seconds)];
```

Esper built-in method.

Patterns: notify if jobcontrol flashlist is not updated within a minute

Variable holding the timeout value minus update interval.

```
create variable int timeout = 50;
```

Creates a separate context partition per each hostname.

```
create context jobctlHostOutdated
context ByHost partition by hostname from jobcontrol,
context Outdated start pattern[a=jobcontrol->
(timer:interval(timeout) and not jobcontrol(hostname=a.hostname))]
end pattern [jobcontrol(hostname=a.hostname)];
```

Ends when the flashlist
with the same hostname arrives.

Starts when the pattern detects
an absence of flashlist.

Accessing context property.

```
context jobctlHostOutdated
select context.ByHost.key1 as hostname,
date(current_timestamp()) as systime,
(current_timestamp()-context.Outdated.a.fetchstamp)/1000.0
as secondsSinceUpdate
from pattern [every timer:interval(10 seconds)];
```

Accessing context starting event.

Match recognize

```
select srcTime, sysTime, fromState, toState
from levelZeroFM_subsys match_recognize(
  partition by SUBSYS
  measures A.STATE as fromState, B.STATE as toState,
         B.timestamp as srcTime, date(B.fetchstamp) as
             sysTime
  after match skip to current row
  pattern (A B)
  define A as A.SUBSYS='DAQ',
         B as B.STATE != A.STATE
);
```

```
select srctime, systime, fromState, toState
  from levelZeroFM_subsys match_recognize(
    partition by SUBSYS
    measures A.STATE as fromState, B.STATE as toState,
           B.timestamp as srctime, date(B.fetchstamp) as
           systime
    after match skip to current row
    pattern (A B)
    define A as A.SUBSYS='DAQ',
           B as B.STATE != A.STATE
  );
```

Match recognize vs patterns

- ▶ Match recognize works with one event type at a time.
- ▶ Match recognize is claimed to be faster and less memory consuming.
- ▶ Match recognize has been proposed for incorporation into SQL standard.

```
select srctime, systime, fromState, toState
from levelZeroFM_subsys match_recognize(
partition by SUBSYS
measures A.STATE as fromState, B.STATE as toState,
        B.timestamp as srctime, date(B.fetchstamp) as
        systime
after match skip to current row
pattern (A B)
define A as A.SUBSYS='DAQ',
        B as B.STATE != A.STATE
);
```

Defines returned values.

Match recognize vs patterns

- ▶ Match recognize works with one event type at a time.
- ▶ Match recognize is claimed to be faster and less memory consuming.
- ▶ Match recognize has been proposed for incorporation into SQL standard.

Match recognize

```
select srctime, systime, fromState, toState
from levelZeroFM_subsys match_recognize(
partition by SUBSYS
measures A.STATE as fromState, B.STATE as toState,
        B.timestamp as srctime, date(B.fetchstamp) as
        systime
after match skip to current row
pattern (A B)
define A as A.SUBSYS='DAQ',
        B as B.STATE != A.STATE
);
```

Defines returned values.

Allows a closing event of one pattern to be the opening event of another.

Match recognize vs patterns

- ▶ Match recognize works with one event type at a time.
- ▶ Match recognize is claimed to be faster and less memory consuming.
- ▶ Match recognize has been proposed for incorporation into SQL standard.

Match recognize

```
select srctime, systime, fromState, toState
from levelZeroFM_subsys match_recognize(
  partition by SUBSYS
  measures A.STATE as fromState, B.STATE as toState,
           B.timestamp as srctime, date(B.fetchstamp) as
           systime
  after match skip to current row
  pattern (A B)
  define A as A.SUBSYS='DAQ',
         B as B.STATE != A.STATE
```

Defines returned values.

Defines temporal pattern:
A directly followed by B.

Allows a closing event of one pattern
to be the opening event of another.

Match recognize vs patterns

- ▶ Match recognize works with one event type at a time.
- ▶ Match recognize is claimed to be faster and less memory consuming.
- ▶ Match recognize has been proposed for incorporation into SQL standard.

Match recognize

```
select srctime, systime, fromState, toState
from levelZeroFM_subsys match_recognize(
partition by SUBSYS
measures A.STATE as fromState, B.STATE as toState,
        B.timestamp as srctime, date(B.fetchstamp) as
        systime
after match skip to current row
pattern (A B)
define A as A.SUBSYS='DAQ',
        B as B.STATE != A.STATE
```

Defines returned values.

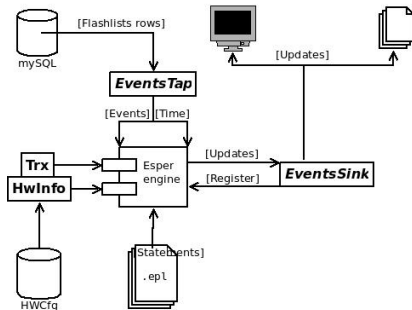
Defines temporal pattern:
A directly followed by B.

Imposes constraints
on events A and B.

Allows a closing event of one pattern
to be the opening event of another.

Match recognize vs patterns

- ▶ Match recognize works with one event type at a time.
- ▶ Match recognize is claimed to be faster and less memory consuming.
- ▶ Match recognize has been proposed for incorporation into SQL standard.



► Events

- unique flashlist rows from November global run
- stored in a mysql database, indexed by reception timestamps (*fetchstamps*)
- flashlists in use:
 - EventProcessorStatus
 - EVM
 - FMMInput
 - FMMStatus
 - frlcontrollerCard
 - frlcontrollerLink
 - gt_cell_lumiseg
 - hostInfo
 - jobcontrol
 - levelZeroFM_static
 - levelZeroFM_subsys
 - StorageManagerPerformance;

► EPL statements stored in files

► EventSinks as statement listeners

► Limitations:

- Non-interactive
- Output repeated for persisting problems.

- ▶ In-run tests:
 - ▶ Message on Run Start giving: run NR, SID, detectors in, FEDs in per detector.
 - ▶ Message on Run Stop giving: run Nr, SID, detectors in, FEDs in per detector, avg L1 rate, avg stream A rate, avg dead time (from trigger LAS), duration.
 - ▶ Message on L1 trigger rate jump $> 10\%$.
 - ▶ Message on subsys going to (ERROR, RUNNING_*, PAUSING, PAUSED, RESUMING)
 - ▶ Message on any state change of DAQ.
 - ▶ Message on jobcontrol flashlist not being updated for 1 minute.
 - ▶ FEDs with dead-time $> 1\%$.
 - ▶ FEDs with backpressure $> 1\%$.
 - ▶ Message on stream A > 500 Hz.
 - ▶ Detection of FEDs in illegal state.
- ▶ In-run test with rate stuck at 0:
 - ▶ Check BX alignment and print message if not aligned. Repeat after 10 seconds.
 - ▶ Check triggers(events) alignment + print message if not aligned. Repeat after 10 seconds.
 - ▶ FEDs stuck in ERROR/OOS/WARNING/BUSY (anything not READY).
 - ▶ List FEDs with backpressure or deadtime.
 - ▶ Check if number of resyncs and the last resync event number is the same for all FEDs.
- ▶ Continuous display:
 - ▶ Stream A rate
 - ▶ L1 rate
 - ▶ Resync rate and number of resyncs
 - ▶ Event processing rate per group (determined by number of processors on machine).
 - ▶ CPU utilization per group (as above)
 - ▶ Deadtime (from trigger LAS)

Swing view screenshot

8u busy ratio	0.003
TRG LAS DT	0.0919
Running	true
system time	Wed Nov 06 23:09:16 CET 2013
#Overperformers	16
RunOngoingContext tmr	1910.0
TRG rate	799.0
AVG rate	782.9
8u events per sec	390.8578210894553
streamA rate	75.22478
Verified FED IDs	518
resync No	0
resync rate	0.0
RunOngoingNestedContext tmr	6150.629
ZeroRateNestedContext tmr	16.0

output	jobctlHostOutdated	output4	output3	context-state-info
RunStart		systemtime: 2013.11.06 20:00:00.408	sid: 225814	runNumber:null
Level 1 TRG rate jump		systemtime: 2013.11.06 20:32:59.379	rate: 852.0	avgrate: 772.2142857142857
Level 1 TRG rate jump		systemtime: 2013.11.06 20:50:22.595	rate: 412.0	avgrate: 772.0714285714286
Level 1 TRG rate jump		systemtime: 2013.11.06 20:50:47.130	rate: 534.0	avgrate: 0.0
FED Dead-time		systemtime: 2013.11.06 20:50:47.747	fedsDt: CSC: [CSC4: (fed#750: 0.163, fed#752: 0.163), CSC-: (fed#754: 0.189, fed#756: 0.189), CSCTF: (fed#760): 0.0]	srcctime: 2013.11.06 19:32:55.2
FED DT gone		systemtime: 2013.11.06 20:50:51.856	feds: CSC: [CSC4: (fed#750, fed#752), CSC-: (fed#754, fed#756), CSCTF: (fed#760)]	srcctime: 2013.11.06 19:50:20.1
Level 1 TRG rate jump		systemtime: 2013.11.06 21:40:49.403	rate: 0.0	avgrate: 777.8125
8X numbers misaligned		systemtime: 2013.11.06 21:41:03.600	bxNumber: 524	fedsNo: 517
8X numbers misaligned		systemtime: 2013.11.06 21:41:03.600	bxNumber: 524	fedsNo: 1
8X numbers misaligned		systemtime: 2013.11.06 21:41:13.709	bxNumber: 524	fedsNo: 517
8X numbers misaligned		systemtime: 2013.11.06 21:41:13.709	bxNumber: 525	fedsNo: 1
systemtime: 2013.11.06 21:41:26.045		msg: 518 trigger numbers scattered into 44 groups of unique values		srcctime: 2013.11.06 20:40:46.8
systemtime: 2013.11.06 21:41:26.045		msg: 518 bx numbers scattered into 37 groups of unique values		srcctime: 2013.11.06 20:41:58.9
Level 1 TRG rate jump		systemtime: 2013.11.06 21:41:26.397	rate: 24.0	avgrate: 0.0
Level 1 TRG rate jump		systemtime: 2013.11.06 21:42:01.937	rate: 120.0	avgrate: 778.0
8X numbers misaligned		systemtime: 2013.11.06 21:42:26.500	bxNumber: 3069	fedsNo: 517
8X numbers misaligned		systemtime: 2013.11.06 21:42:26.500	bxNumber: 3070	fedsNo: 1
Level 1 TRG rate jump		systemtime: 2013.11.06 21:42:31.173	rate: 790.0	avgrate: 0.0
Subsystem changed state		systemtime: 2013.11.06 22:27:14.630	subsys: TRG	fromState:Running
Subsystem changed state		systemtime: 2013.11.06 22:27:18.852	subsys: TRG	fromState:Pausing
DAQ changed state		systemtime: 2013.11.06 22:27:18.852	fromState:Running	toState: Stopping
Run Stop		systemtime: 2013.11.06 22:27:18.852	sid: 225814	runNumber:216311
Level 1 TRG rate jump		systemtime: 2013.11.06 22:27:20.976	rate: 0.0	avgrate: 782.6666666666666
DAQ changed state		systemtime: 2013.11.06 22:27:31.446	fromState:Stopping	toState: Configured
DAQ changed state		systemtime: 2013.11.06 22:32:10.829	fromState:Configured	toState: Standby
DAQ changed state		systemtime: 2013.11.06 22:32:16.094	fromState:Standby	toState: Initializing
DAQ changed state		systemtime: 2013.11.06 22:32:35.090	fromState:Initializing	toState: Halted
DAQ changed state		systemtime: 2013.11.06 22:32:58.575	fromState:Halted	toState: Initialized
DAQ changed state		systemtime: 2013.11.06 22:33:01.174	fromState:Initialized	toState: Halted
DAQ changed state		systemtime: 2013.11.06 22:35:27.170	fromState:Halted	toState: Configuring

DaqDoctor's output

2013-11-08 07:10:35	NO BEAM	FED BX numbers are not aligned. 1 different bc numbers encountered
2013-11-08 07:10:35	NO BEAM	Most (475) had the bc set to 2454 in the last event seen. 28 frls saw the last bc at 2442: HBHEb (711), HBHEb (706), HBHEb (709), HBHEb (708), HBHEb (707), HBHEb (710), HBHEc (712), HBHEc (715), HBHEc (716), HBHEc (714), HBHEc (717), HBHEc (713), HO (727), HO (728), HO (730), HO (724), HO (731), HO (729), HO (725), HO (726), GT (812), GT (813), HBHEa (701), HBHEa (703), HBHEa (705), HBHEa (702), HBHEa (704), HBHEa (700) 1 frls saw the last bc at 2446: DTTF (780) 1 frls saw the last bc at 2443: SCAL (735)

Prototype's output: BX numbers misaligned

sysTime	bxn	fedsNo	feds
2013.11.08 07:10:39.646	2454	475	[REMOVED]
2013.11.08 07:10:39.646	3490	28	HCal: [HBHEa: (700, 701, 702, 703, 704, 705), HBHEb: (706, 707, 708, 709, 710, 711), HBHEc: (712, 713, 714, 715, 716, 717), HO: (724, 725, 726, 727, 728, 729, 730, 731)], TRG: [GT: (812, 813)]
2013.11.08 07:10:39.646	3491	1	SCAL: [SCAL: (735)]
2013.11.08 07:10:39.646	3494	1	DT: [DTTF: (780)]

Tested on a set of 111M events recorded between 2013-11-06 20:00:00 and 2013-11-08 20:00:00.

- ▶ Engine clock, controlled by incoming data, progresses 42 times as fast as system clock (Intel i7 CPU 4x3.40GHz).
- ▶ This metric reaches 44 with no statements or no data sent to the engine.
- ▶ It reaches 34000 if, additionally, nothing but timestamps is being pulled from DB.

Pros:

- ▶ An out-of-the-box event processing framework.
- ▶ Extensibility and ease of embedding.
- ▶ Separation of event-processing logic and Java code.
- ▶ Satisfactory documentation.

Cons:

- ▶ EPL
 - ▶ sometimes too rigid and conter-intuitive:
 - ▶ Variables can only be set using triggering event in ON-SET statements,
 - ▶ pattern-triggered queries cannot select from unnamed windows (`unidirectional join` as a workaround)
 - ▶ incomplete: e.g. forbids using `rstream` in patterns
 - ▶ evolving:
 - ▶ bugs fixed with each version (prototype crashes on 4.9)
 - ▶ `group by` in subqueries since 4.11
 - ▶ subqueries within `having` since 4.10
 - ▶ steep learning curve
 - ▶ cumbersome data transformations
- ▶ Hard to debug.
- ▶ Lack of community, books, thorough tutorials.

Questions, suggestions?

Thank you!

Thank you!

Thank you!

```
create objectarray schema ResyncStream as (resyncRate double,  
    resyncNo long);  
create variable String desiredContext = "";  
create variable Integer desiredSlotNumber = -1;  
  
on FrlCtlLnk(fedSrcId=812) as ctl set desiredContext = ctl.context,  
    desiredSlotNumber= ctl.slotNumber;  
  
insert into ResyncStream  
select resyncRate, resyncNo from frlcontrollerCard match_recognize(  
    partition by context, slotNumber  
    measures 1000*(B.myrinetResync - A.myrinetResync)/(B.timestamp.  
        getTime()-A.timestamp.getTime()) as resyncRate, B.  
        myrinetResync as resyncNo  
    after match skip to current row  
    pattern (A B)  
    define A as A.context = desiredContext and A.slotNumber=  
        desiredSlotNumber  
);
```

```
create objectarray schema RunAvgRate as (rate double);
```

```
create objectarray schema RunAvgRate as (rate double);  
  
select avg(rate) as rate from L1RatesStream;
```

```
create objectarray schema RunAvgRate as (rate double);
```

```
select avg(rate) as rate from L1RatesStream;
```

Using contexts to reinit results.

Expression with a subquery.

```
create expression runAvgRate{  
  (select rate from RunAvgRate.win:length(1))};  
create expression runNumber{  
  (select RUN_NUMBER from levelZeroFM_static.std:unique(SID) where  
    SID=sid)};
```

```
create objectarray schema RunAvgRate as (rate double);
```

```
select avg(rate) as rate from L1RatesStream;
```

Using contexts to reinit results.

Expression with a subquery.

```
create expression runAvgRate{  
  (select rate from RunAvgRate.win:length(1))};  
create expression runNumber{  
  (select RUN_NUMBER from levelZeroFM_static.std:unique(SID) where  
   SID=sid)};
```

Stream with a combination of views.

```
runAvgRate() as avgRunRate, formatMs(runDuration()) as  
  runDuration ...  
from pattern [every p=RunStop];
```

```
create objectarray schema RunAvgRate as (rate double);
```

```
select avg(rate) as rate from L1RatesStream;
```

Using contexts to reinit results.

Expression with a subquery.

```
create expression runAvgRate{  
  (select rate from RunAvgRate.win:length(1))};  
create expression runNumber{  
  (select RUN_NUMBER from levelZeroFM_static.std:unique(SID) where  
   SID=sid)};
```

Stream with a combination of views.

```
runAvgRate() as avgRunRate, formatMs(runDuration()) as  
  runDuration ...  
from pattern [every p=RunStop];
```



```
create objectarray schema RunAvgRate as (rate double);
```

```
select avg(rate) as rate from L1RatesStream;
```

Using contexts to reinit results.

Expression with a subquery.

```
create expression runAvgRate{  
  (select rate from RunAvgRate.win:length(1))};  
create expression runNumber{  
  (select RUN_NUMBER from levelZeroFM_static.std:unique(SID) where  
   SID=sid)};
```

Stream with a combination of views.

```
runAvgRate() as avgRunRate, formatMs(runDuration()) as  
  runDuration ...  
from pattern [every p=RunStop];
```

```
create objectarray schema RunAvgRate as (rate double);
```

```
select avg(rate) as rate from L1RatesStream;
```

Using contexts to reinit results.

Expression with a subquery.

```
create expression runAvgRate{  
  (select rate from RunAvgRate.win:length(1))};  
create expression runNumber{  
  (select RUN_NUMBER from levelZeroFM_static.std:unique(SID) where  
   SID=sid)};
```

Stream with a combination of views.

```
runAvgRate() as avgRunRate, formatMs(runDuration()) as  
  runDuration ...  
from pattern [every p=RunStop];
```

```
create objectarray schema RunAvgRate as (rate double);
```

```
select avg(rate) as rate from L1RatesStream;
```

Using contexts to reinit results.

```
create expression runAvgRate{  
  (select rate from RunAvgRate.win:length(1))};  
create expression runNumber{  
  (select RUN_NUMBER from levelZeroFM_static.std:unique(SID) where  
    SID=sid)};
```

Expression with a subquery.

Expression invokations.

Stream with a combination of views.

```
runAvgRate() as avgRunRate, formatMs(runDuration()) as  
  runDuration ...  
from pattern [every p=RunStop];
```

```
create objectarray schema RunAvgRate as (rate double);
```

```
select avg(rate) as rate from L1RatesStream;
```

Using contexts to reinit results.

```
create expression runAvgRate{  
  (select rate from RunAvgRate.win:length(1))};  
create expression runNumber{  
  (select RUN_NUMBER from levelZeroFM_static.std:unique(SID) where  
    SID=sid)};
```

Expression with a subquery.

Expression invocations.

Stream with a combination of views.

```
runAvgRate() as avgRunRate, formatMs(runDuration()) as  
  runDuration ...  
from pattern [every p=RunStop];
```

Using a registered single-row function.

```
create objectarray schema RunAvgRate as (rate double);
```

```
select avg(rate) as rate from L1RatesStream;
```

Using contexts to reinit results.

```
create expression runAvgRate{  
  (select rate from RunAvgRate.win:length(1))};  
create expression runNumber{  
  (select RUN_NUMBER from levelZeroFM_static.std:unique(SID) where  
    SID=sid)};
```

Expression with a subquery.

Expression invocations.

Stream with a combination of views.

```
runAvgRate() as avgRunRate, formatMs(runDuration()) as  
  runDuration ...  
from pattern [every p=RunStop];
```

Using a registered single-row function.

Extensibility

Esper allows users to provide custom aggregation functions, single row functions, annotations, views and others.

Enumeration methods & output clause

```
context L1ZeroRate
select
  date(max(b.timestamp)) as srctime,
  date(current_timestamp()) as systime,
  fedsInfoString(window(b.*).toMap(
    k=>k.fedSrcId, v=>format(v.bpFraction))) as feds
from
  FrlBackpressure(timestamp is not null) as b,
  pattern[every FrlBackpressure->(timer:interval(1 msec)
    and not FrlBackpressure)] unidirectional
output first every 5 seconds;
```

```
context L1ZeroRate
select
  date(max(b.timestamp)) as srctime,
  date(current_timestamp()) as systime,
  fedsInfoString(window(b.*).toMap(
    k=>k.fedSrcId, v=>format(v.bpFraction))) as feds
from
  FrlBackpressure(timestamp is not null) as b,
  pattern[every FrlBackpressure->(timer:interval(1 msec)
    and not FrlBackpressure)] unidirectional
output first every 5 seconds;
```

Enumeration methods & output clause

Using a helper method `fedsInfoString(Map<Long,String>)`.

```
context L1ZeroRate
select
  date(max(b.timestamp)) as srctime,
  date(current_timestamp()) as systime,
  fedsInfoString(window(b.*).toMap(
    k=>k.fedSrcId, v=>format(v.bpFraction))) as feds
from
  FrlBackpressure(timestamp is not null) as b,
  pattern[every FrlBackpressure->(timer:interval(1 msec)
    and not FrlBackpressure)] unidirectional
output first every 5 seconds;
```


Enumeration methods & output clause

window() aggregation function.

```
context L1ZeroRate
select
  date(max(b.timestamp)) as srctime,
  date(current_timestamp()) as systime,
  fedsInfoString(window(b.*).toMap(
    k=>k.fedSrcId, v=>format(v.bpFraction))) as feds
from
  FrlBackpressure(timestamp is not null) as b,
  pattern[every FrlBackpressure->(timer:interval(1 msec)
    and not FrlBackpressure)] unidirectional
output first every 5 seconds;
```

Enumeration methods & output clause

window() aggregation function.

toMap() enumeration function.

```
context L1ZeroRate
select
  date(max(b.timestamp)) as srctime,
  date(current_timestamp()) as systime,
  fedsInfoString(window(b.*).toMap(
    k=>k.fedSrcId, v=>format(v.bpFraction))) as feds
from
  FrlBackpressure(timestamp is not null) as b,
  pattern[every FrlBackpressure->(timer:interval(1 msec)
    and not FrlBackpressure)] unidirectional
output first every 5 seconds;
```

Enumeration methods & output clause

window() aggregation function.

toMap() enumeration function.

```
context L1ZeroRate
select
  date(max(b.timestamp)) as srctime,
  date(current_timestamp()) as systime,
  fedsInfoString(window(b.*).toMap(
    k=>k.fedSrcId, v=>format(v.bpFraction))) as feds
from
  FrlBackpressure(timestamp is not null) as b,
  pattern[every FrlBackpressure->(timer:interval(1 msec)
    and not FrlBackpressure)] unidirectional
output first every 5 seconds;
```

Enumeration methods: performing common tasks on collections of events

Some, like the aggregate() method, accept lambda expressions:

```
window(b.*).aggregate("", (result, row)=>result||" "||row.fedSrcId)
```

Enumeration methods & output clause

```
context L1ZeroRate
select
  date(max(b.timestamp)) as srctime,
  date(current_timestamp()) as systime,
  fedsInfoString(window(b.*).toMap(
    k=>k.fedSrcId, v=>format(v.bpFraction))) as feds
from
  FrlBackpressure(timestamp is not null) as b,
  pattern[every FrlBackpressure->(timer:interval(1 msec)
    and not FrlBackpressure)] unidirectional
output first every 5 seconds;
```

Annotations:

- window() aggregation function.
- toMap() enumeration function.
- Using a helper method format(double).

Enumeration methods: performing common tasks on collections of events

Some, like the aggregate() method, accept lambda expressions:

```
window(b.*).aggregate("", (result, row)=>result||" "||row.fedSrcId)
```

Enumeration methods & output clause

```
context L1ZeroRate
```

```
select
```

```
  date(max(b.timestamp)) as srctime,
```

```
  date(current_timestamp()) as systime,
```

```
  fedsInfoString(window(b.*).toMap(
```

```
    k=>k.fedSrcId, v=>format(v.bpFraction))) as feds
```

```
from
```

```
  FrlBackpressure(timestamp is not null) as b,
```

```
  pattern[every FrlBackpressure->(timer:interval(1 msec)
```

```
    and not FrlBackpressure)] unidirectional
```

```
output first every 5 seconds;
```

window() aggregation function.

toMap() enumeration function.

Pattern triggered selection.

Enumeration methods: performing common tasks on collections of events

Some, like the aggregate() method, accept lambda expressions:

```
window(b.*).aggregate("", (result, row)=>result||" "||row.fedSrcId)
```

Enumeration methods & output clause

```
context L1ZeroRate
select
  date(max(b.timestamp)) as srctime,
  date(current_timestamp()) as systime,
  fedsInfoString(window(b.*).toMap(
    k=>k.fedSrcId, v=>format(v.bpFraction))) as feds
from
  FrlBackpressure(timestamp is not null) as b,
  pattern[every FrlBackpressure->(timer:interval(1 msec)
    and not FrlBackpressure)] unidirectional
output first every 5 seconds;
```

window() aggregation function.

toMap() enumeration function.

Output stabilizing clause.

Enumeration methods: performing common tasks on collections of events

Some, like the aggregate() method, accept lambda expressions:

```
window(b.*).aggregate("", (result, row)=>result||" "||row.fedSrcId)
```

Enumeration methods & output clause

```
context L1ZeroRate
select
  date(max(b.timestamp)) as srctime,
  date(current_timestamp()) as systime,
  fedsInfoString(window(b.*).toMap(
    k=>k.fedSrcId, v=>format(v.bpFraction))) as feds
from
  FrlBackpressure(timestamp is not null) as b,
  pattern[every FrlBackpressure->(timer:interval(1 msec)
    and not FrlBackpressure)] unidirectional
output first every 5 seconds;
```

window() aggregation function.

toMap() enumeration function.

Output stabilizing clause.

Enumeration methods: performing common tasks on collections of events

Some, like the aggregate() method, accept lambda expressions:

```
window(b.*).aggregate("", (result, row)=>result||" "||row.fedSrcId)
```

Output clause: suppressing output and controlling its rate

The suppressing logic allows reading and setting variables, using crontab notation and more.